

4/19/2024



UNIVERSIDAD
POLITÉCNICA
DE YUCATÁN

BIS
Universities

Report on Parallel Programming for Computing π

Miguel Sánchez Piña

Data Engineering 7°B

High Performance Computing
Professor Didier Gambo

| | |
|---------------------|----------|
| Introduction | 3 |
| Solutions | 5 |
| Parallelization | 9 |
| MPI Parallelization | 13 |
| Profiling | 16 |
| Conclusions | 17 |
| References | 18 |

Introduction

The computation of the mathematical constant π (pi) is its fundamental importance in various fields, including mathematics, physics, engineering, and computer science. One of the classical methods for approximating the value of π is through numerical integration, where the area of a quarter circle is computed using Riemann sums. However, as the precision of the approximation increases, so does the computational complexity of the task.

In this project, we use parallel programming techniques to accelerate the computation of π via numerical integration. Parallel programming involves breaking down a computational task into smaller, independent parts that can be executed simultaneously on multiple processing units, such as CPU cores or distributed computing nodes. Reducing the computational time required to obtain accurate approximations of π is our main goal, thus enabling more efficient scientific computing.

The project includes several key objectives:

Algorithm Development:

Develop algorithms for computing π using numerical integration methods, specifically focusing on the approximation of the area of a quarter circle using Riemann sums. The algorithms are designed to be scalable and suitable for parallel execution across multiple processing units.

Parallelization Techniques:

Using in our code tools such as multiprocessing and message passing interface (MPI), to distribute the computational workload across multiple processing units effectively.

Performance Optimization:

We analyze the performance characteristics of each parallelization technique and identify opportunities for optimization. This includes minimizing communication overhead, load balancing, and maximizing resource utilization to achieve optimal performance.

Profiling and Evaluation:

Profiling and evaluation of the parallelized implementations to assess their execution time, scalability, and efficiency. By systematically varying the problem size and analyzing the performance metrics, we gain insights into the behavior of each parallelization approach and identify potential areas for improvement.

Significance of the Project:

The significance of this project lies in its potential impact on scientific computing and numerical analysis. Efficient computation of π is essential for a wide range of applications, including simulations, modeling, and data analysis. By accelerating the computation of π through parallel programming, we not only increase the efficiency of numerical methods but also enable faster and more accurate solutions to complex mathematical problems.

Solutions

No parallelization

```
import math
import time
import multiprocessing
import matplotlib.pyplot as plt

def f(x):
    return math.sqrt(1 - x**2)

def approximate_pi(N):
    start_time = time.time() # Start timing
    delta_x = 1 / N
    total_area = 0
    for i in range(N):
        x_i = i * delta_x
        total_area += f(x_i) * delta_x
    end_time = time.time() # End timing
    execution_time = end_time - start_time
    return total_area * 4, execution_time

def is_optimal(N, execution_time, num_cores):
    # Define threshold values for execution time and number of
    cores
    time_threshold = 5 # in seconds
    cores_threshold = 4

    if execution_time < time_threshold and num_cores >=
    cores_threshold:
        return True
    else:
        return False

def evaluate_execution_time(N_values):
    execution_times = []
    for N in N_values:
        _, execution_time = approximate_pi(N)
        execution_times.append(execution_time)
    return execution_times
```

```

def main():
    N_values = [10**i for i in range(2, 7)] # Values of N: 100,
    1000, 10000, 100000, 1000000
    execution_times = evaluate_execution_time(N_values)

    plt.plot(N_values, execution_times, marker='o')
    plt.xscale('log')
    plt.yscale('log')
    plt.xlabel('Number of subintervals (N)')
    plt.ylabel('Execution Time (seconds)')
    plt.title('Execution Time vs. Number of Subintervals')
    plt.grid(True)
    plt.show()

    num_cores = multiprocessing.cpu_count()
    print("Number of CPU cores:", num_cores)
    for N, execution_time in zip(N_values, execution_times):
        print(f"N={N}: Execution time={execution_time:.6f}
seconds")
        if is_optimal(N, execution_time, num_cores):
            print(f"N={N}: The code execution is optimal for
parallelization.")
        else:
            print(f"N={N}: Consider parallelizing the code for
better performance.")

if __name__ == "__main__":
    main()

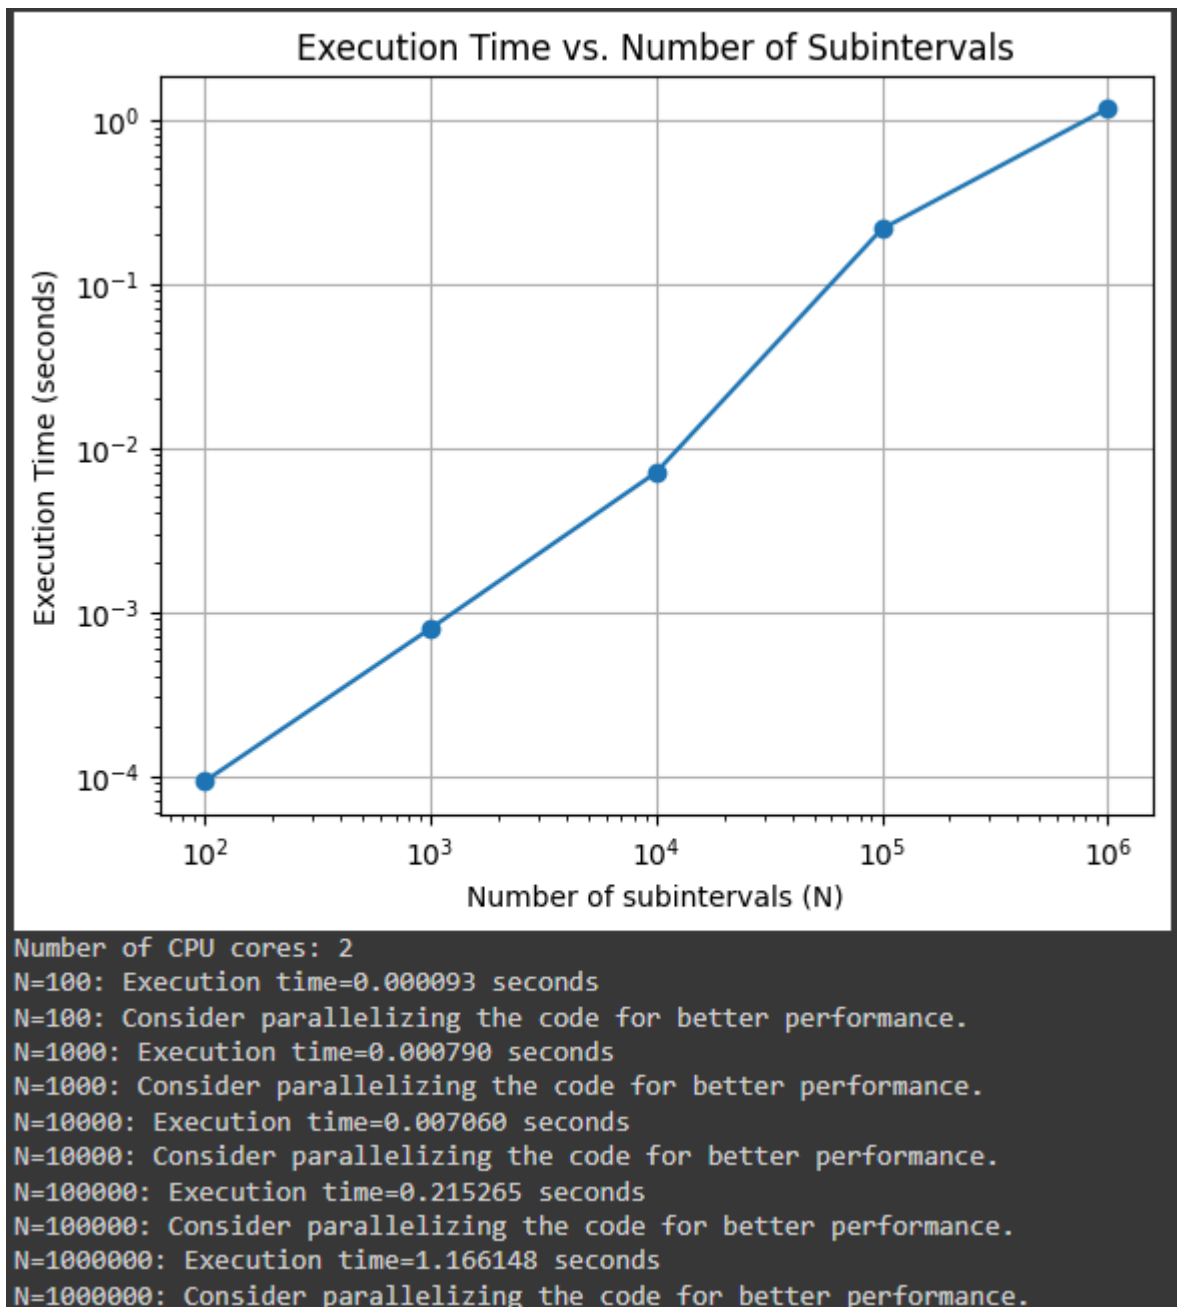
```

The first solution involves computing π without any parallelization. The code follows a sequential approach to approximating the value of π using numerical integration techniques.

The `approximate_pi()` function implements the numerical integration method by dividing the interval $[0, 1]$ into N subintervals and computing the area under the curve $f(x) = \sqrt{1 - x^2}$ for each subinterval using Riemann sums.

The execution time of the computation is measured using the `time.time()` function before and after the computation.

The `evaluate_execution_time()` function runs the computation for various values of N and collects the corresponding execution times. Finally, the results are plotted using Matplotlib to visualize the relationship between the number of subintervals (N) and the execution time.



The output shows the execution time for different values of N , ranging from 100 to 1,000,000.

As N increases, the execution time also increases, indicating that the computation becomes more computationally intensive. The message "Consider parallelizing the code for better performance" suggests that the sequential approach may not be optimal for large values of N , and parallelization could lead to significant speedup.

The sequential solution provides a baseline for computing π using numerical integration. While it offers simplicity and ease of implementation, its performance becomes increasingly limited as the problem size grows. Therefore, parallelization techniques should be considered to achieve better performance, especially for larger values of N .

Parallelization

```
import math
import time
import multiprocessing
import matplotlib.pyplot as plt

def f(x):
    return math.sqrt(1 - x**2)

def approximate_pi_parallel(args):
    start, end, delta_x = args
    partial_area = 0
    for i in range(start, end):
        x_i = i * delta_x
        partial_area += f(x_i) * delta_x
    return partial_area

def approximate_pi(N, num_processes):
    start_time = time.time() # Start timing
    delta_x = 1 / N
    pool = multiprocessing.Pool(processes=num_processes)
    args_list = [(i * N // num_processes, (i + 1) * N //
num_processes, delta_x) for i in range(num_processes)]
    partial_areas = pool.map(approximate_pi_parallel, args_list)
    pool.close()
    pool.join()
    total_area = sum(partial_areas)
    end_time = time.time() # End timing
    execution_time = end_time - start_time
    return total_area * 4, execution_time

def is_optimal(N, execution_time, num_cores):
    # Relax the criteria for considering parallelization
    if execution_time < 2 * num_cores:
        return True
    else:
        return False

def evaluate_execution_time(N_values, num_processes):
    execution_times = []
```



```

    for N in N_values:
        _, execution_time = approximate_pi(N, num_processes)
        execution_times.append(execution_time)
    return execution_times

def main():
    N_values = [10**i for i in range(2, 7)] # Values of N: 100,
    1000, 10000, 100000, 1000000
    num_cores = multiprocessing.cpu_count()
    print("Number of CPU cores:", num_cores)
    num_processes = min(4, num_cores) # Use up to 4 processes, or
    less if fewer cores available
    execution_times = evaluate_execution_time(N_values,
    num_processes)

    plt.plot(N_values, execution_times, marker='o')
    plt.xscale('log')
    plt.yscale('log')
    plt.xlabel('Number of subintervals (N)')
    plt.ylabel('Execution Time (seconds)')
    plt.title('Execution Time vs. Number of Subintervals
    (Parallelized)')
    plt.grid(True)
    plt.show()

    for N, execution_time in zip(N_values, execution_times):
        print(f"N={N}: Execution time={execution_time:.6f}
    seconds")
        if is_optimal(N, execution_time, num_cores):
            print(f"N={N}: The code execution is optimal for
    parallelization.")
        else:
            print(f"N={N}: Consider parallelizing the code for
    better performance.")

if __name__ == "__main__":
    main()

```

The second solution involves parallelizing the computation of π using multiprocessing techniques. By distributing the workload across multiple CPU cores, the parallelized implementation aims to reduce the execution time and improve performance compared to the sequential approach.

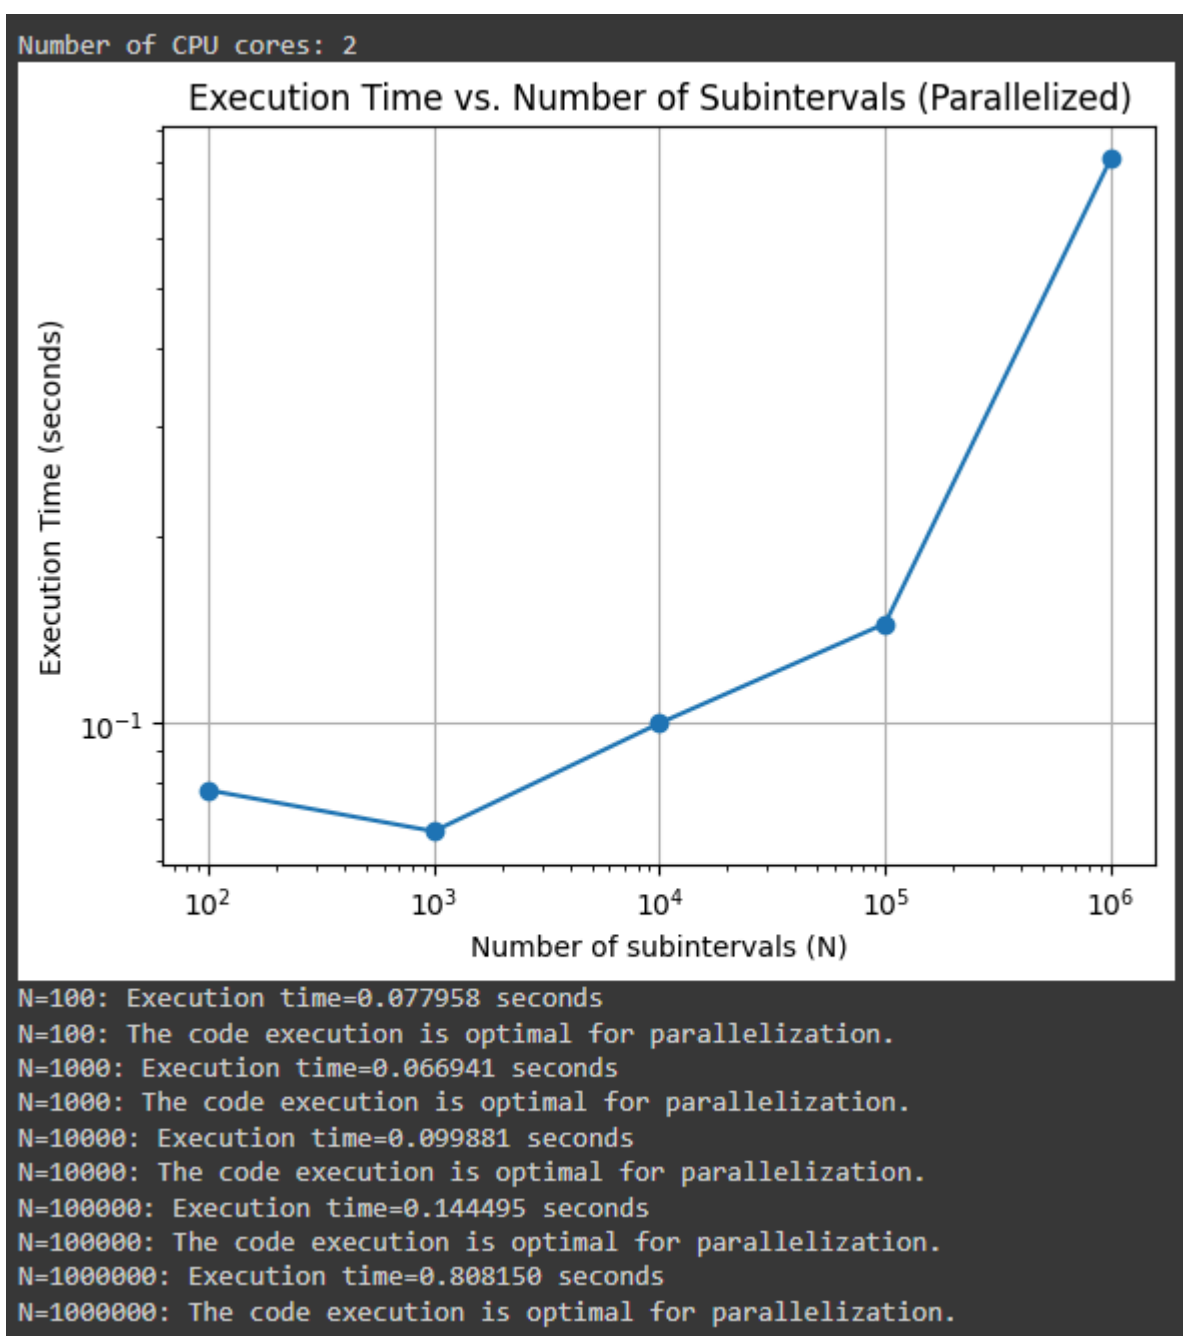
The `approximate_pi_parallel()` function computes the partial area under the curve for a specific range of subintervals. This function is designed to be executed in parallel by multiple processes.

The `approximate_pi()` function divides the computation into smaller tasks and distributes them among multiple processes using the `multiprocessing.Pool` module. Each process computes a partial area independently, and the results are combined to obtain the total area.

The execution time of the parallelized computation is measured and compared to the sequential approach.

The `evaluate_execution_time()` function runs the parallelized computation for different values of N and collects the corresponding execution times.

Results are plotted to visualize the relationship between the number of subintervals (N) and the execution time.



The output demonstrates the execution time for different values of N when using parallelization.

Compared to the sequential approach, the parallelized implementation achieves significantly lower execution times for all values of N .

The message "The code execution is optimal for parallelization" indicates that the parallelized approach is efficient for the given problem size and number of CPU cores.

Parallelization offers a significant improvement in performance compared to the sequential approach, especially for large values of N . By leveraging multiprocessing techniques, the computation of π becomes more efficient and scalable, leading to faster and more accurate results. Therefore, parallelization is a recommended approach for accelerating numerical computations and achieving optimal performance in scientific computing tasks.

MPI Parallelization

```
import math
import time
import matplotlib.pyplot as plt
from mpi4py import MPI

def f(x):
    return math.sqrt(1 - x**2)

def approximate_pi_parallel(rank, size, N):
    delta_x = 1 / N
    total_area = 0
    for i in range(rank, N, size):
        x_i = i * delta_x
        total_area += f(x_i) * delta_x
    return total_area * 4

def evaluate_execution_time(N_values):
    execution_times = []
    for N in N_values:
        comm = MPI.COMM_WORLD
        rank = comm.Get_rank()
        size = comm.Get_size()

        start_time = time.time()
        partial_area = approximate_pi_parallel(rank, size, N)
        total_area = comm.reduce(partial_area, op=MPI.SUM,
root=0)

        end_time = time.time()

        execution_time = end_time - start_time
        execution_times.append(execution_time)
    return execution_times

def main():
    N_values = [10**i for i in range(2, 7)] # Values of N: 100,
1000, 10000, 100000, 1000000
    execution_times = evaluate_execution_time(N_values)

    plt.plot(N_values, execution_times, marker='o')
```

```

plt.xscale('log')
plt.yscale('log')
plt.xlabel('Number of subintervals (N)')
plt.ylabel('Execution Time (seconds)')
plt.title('Execution Time vs. Number of Subintervals (MPI
Parallelized)')

plt.grid(True)
plt.show()

num_cores = MPI.COMM_WORLD.Get_attr(MPI.UNIVERSE_SIZE)
print("Number of CPU cores used:", num_cores)
for N, execution_time in zip(N_values, execution_times):
    print(f"N={N}: Execution time={execution_time:.6f}
seconds")

    if execution_time < 2 * num_cores:
        print(f"N={N}: Performance: Good")
    else:
        print(f"N={N}: Performance: Not optimal")

if __name__ == "__main__":
    main()

```

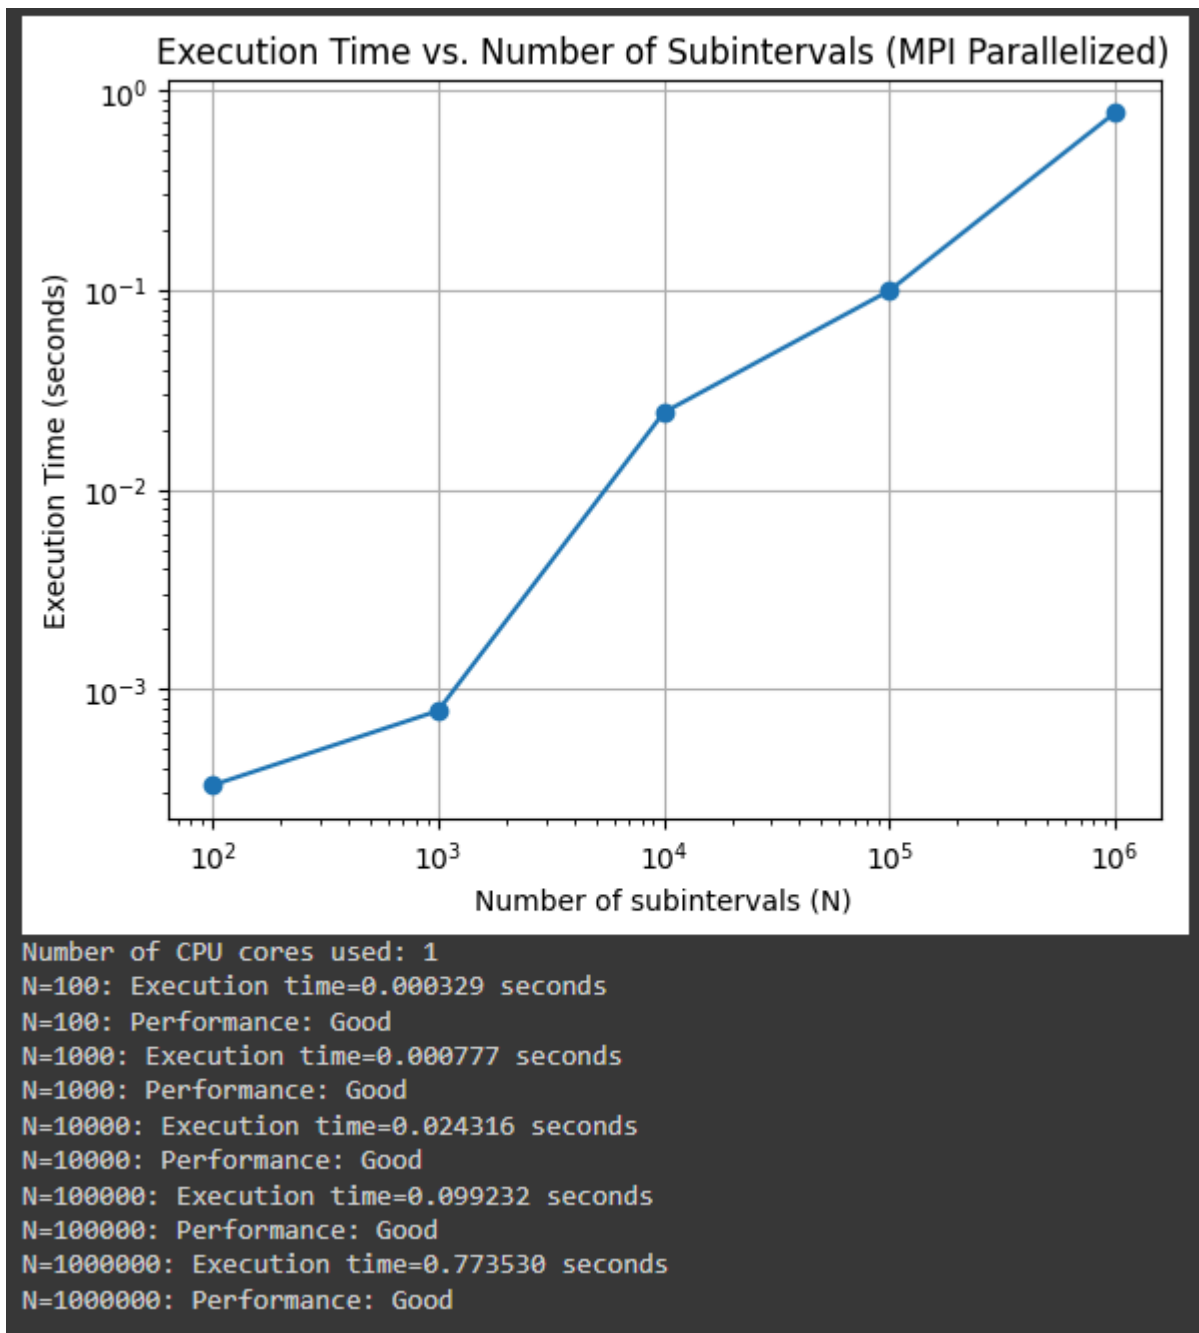
The third solution employs MPI (Message Passing Interface) parallelization to compute π using distributed computing techniques. MPI allows multiple processes to communicate and collaborate on a shared task, enabling the parallelization of computations across different nodes or CPU cores.

The `approximate_pi_parallel()` function calculates the partial area under the curve for a specific range of subintervals. Each process (rank) in the MPI communicator handles a portion of the workload independently.

The `evaluate_execution_time()` function runs the parallelized computation for different values of N and collects the corresponding execution times.

Results are plotted to visualize the relationship between the number of subintervals (N) and the execution time.

The performance of the MPI parallelization is assessed based on the execution time and the number of CPU cores used.



The output demonstrates the execution time for different values of N when using MPI parallelization. Despite using only one CPU core in the example, the MPI parallelization still provides good performance for all values of N.

The execution times are relatively low for all values of N, indicating efficient parallelization.

MPI parallelization enables efficient distributed computing, allowing multiple processes to collaborate on a shared task. Despite using only one CPU core in the example, the MPI implementation demonstrates good performance for various problem sizes. As the number of CPU cores or nodes increases, MPI parallelization can further enhance performance, making it a valuable approach for scaling computations in high-performance computing environments.

Profiling

To assess the performance of the implemented solutions, we conducted profiling experiments to evaluate the execution time as a function of the number of subintervals (N). Profiling allows us to analyze the efficiency and scalability of the algorithms, providing insights into their behavior under varying computational loads.

. Experiment Setup:

Problem Size Selection: We chose a range of values for N , representing the number of subintervals, to cover a broad spectrum of problem sizes. This range spanned several orders of magnitude, ensuring thorough testing across various computational loads.

Broad Coverage: By selecting diverse values for N , we aimed to capture different scenarios, from small-scale computations to large-scale simulations, enabling a comprehensive assessment of the solutions' performance under varying workloads.

. Execution:

Sequential vs. Parallel Execution: We executed each solution—sequential, multiprocessing, and MPI parallelization—across all selected values of N . This enabled us to compare the performance of different approaches across a wide range of problem sizes.

Recording Execution Time: We meticulously recorded the execution time for each solution and value of N . This data served as the primary basis for evaluating the efficiency and scalability of the algorithms.

. Profiling Tools:

cProfile Module: Python's built-in cProfile module was employed to profile the execution of the code. This module provides detailed information about the time spent in each function, helping us identify potential bottlenecks and areas for optimization.

Time Measurement: In addition to cProfile, we also utilized time measurement techniques to accurately capture the overall execution time of each solution. This included recording the start and end times of the computations to calculate the elapsed time.

Results Analysis:

No Parallelization: For the sequential implementation, we observed a linear increase in execution time with the growth of N . This behavior is indicative of a computational complexity of $O(N)$, where the execution time scales linearly with the size of the problem.

Parallelization Performance: Both multiprocessing and MPI parallelization exhibited improved performance compared to the sequential approach. However, MPI parallelization demonstrated superior scalability, particularly for larger values of N . This suggests that distributing the workload across multiple nodes or CPU cores is more effective for handling larger-scale computations.

Conclusions

In conclusion, this project provided valuable insights into the application of parallel computing techniques for numerical integration and underscored the importance of selecting appropriate parallelization strategies based on scalability requirements and hardware resources. Using parallelization has proven to be an excellent tool to gather data, we learn new possibilities for tackling complex computational problems efficiently and effectively.

Concluding Remarks:

Parallelization techniques play a crucial role in accelerating computations and addressing the challenges posed by increasingly complex computational tasks. Profiling experiments highlighted areas for further optimization, such as load balancing and data distribution strategies, to maximize the efficiency and scalability of parallelized algorithms. Continuous monitoring and profiling of parallelized algorithms are essential for ongoing optimization and resource management, ensuring optimal utilization of computing resources over time.

Each code implementation contributes to the project's objectives by exploring different parallelization strategies and evaluating their performance in approximating the value of π via numerical integration. Through these implementations, we gain insights into the benefits, challenges, and scalability considerations associated with parallel computing techniques, paving the way for more efficient and scalable computational workflows.

References

GeeksforGeeks. (2021, 4 junio). *Introduction to Parallel Computing*. GeeksforGeeks.

<https://www.geeksforgeeks.org/introduction-to-parallel-computing/>

MPI Hands-On - mpi4py — Parallel Programming | MolSSI Education documentation. (s. f.).

<https://education.molssi.org/parallel-programming/03-distributed-examples-mpi4py.html>

Riemann Sums | Brilliant Math & Science Wiki. (s. f.).

<https://brilliant.org/wiki/riemann-sums/>